

**Department of Electronic and
Telecommunication Engineering
University of Moratuwa**

**EN3160 - Image Processing and Machine
Vision**



Assignment 01

210668P Vidmal H.V.P.

1 Question 1: Intensity Transformation

```

1 t1 = np.linspace(0, 50, num=51).astype('uint8')
2 t2 = np.linspace(100, 255, num=100).astype('uint8')
3 t3 = np.linspace(150, 255, num=105).astype('uint8')
4 # Concatenate the arrays
5 t = np.concatenate((t1, t2, t3), axis=0).astype('uint8')
6 print(t.shape)
7 # Apply the transformation
8 f = t[image]

```



Discussion: This intensity transformation enhances mid-range gray values while preserving extreme light and dark areas. Its key feature is abrupt transitions in the mapping function, creating sharp contrast changes in mid-tones. These discontinuities produce a visually striking effect, increasing contrast in specific regions while maintaining the original image's overall structure.

2 Question 2: Accentuating White and Gray Matter

```

1 #white matter
2 mu = 150
3 sigma = 25
4 x = np.linspace(0, 255, 256)
5 t1 = 255 * np.exp(-((x - mu)*2) / (2 * sigma*2))
6 #gray matter
7 mu = 210
8 sigma = 25
9 x = np.linspace(0, 255, 256)
10 t2 = 255 * np.exp(-((x - mu)*2) / (2 * sigma*2))

```

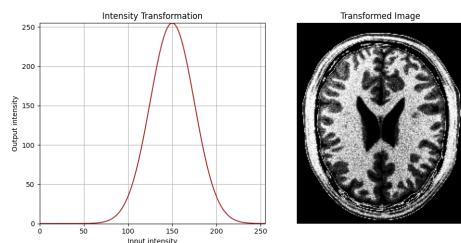


Figure 1: White Matter

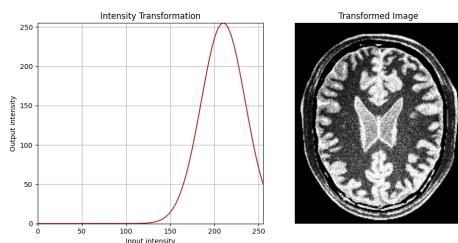


Figure 2: Gray Matter

Discussion The Gaussian transformations effectively enhanced specific brain tissues in the MRI. Centering at intensity 150 highlighted white matter, while centering

at 210 emphasized grey matter. These results demonstrate the method's ability to selectively enhance different anatomical features. However, the original image's inherent noise remained unaffected by the transformations, persisting in both enhanced images.

3 Question 3: Gamma Correction in L*a*b* Color Space

```
1 L,a,b = cv2.split(shadow_image)
2
3 gamma = 0.6
4 t = np.array([(i/255.0)**(gamma)*255 for i in np.arange(0, 256)
5             ]).astype('uint8')
6 new_shadow_image = cv2.LUT(L, t)
7
8 merge_image = cv2.merge([new_shadow_image, a, b])
9 after_new_shadow_image= cv2.cvtColor(merge_image, cv2.
10 COLOR_LAB2RGB)
```



Figure 3: WOriginal and gamma corrected output

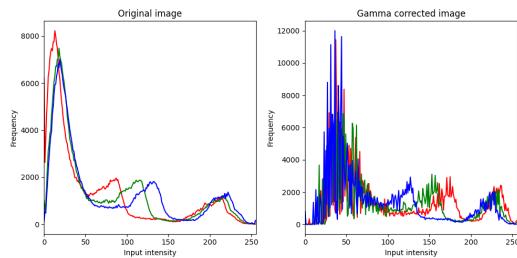


Figure 4: Histograms

Discussion Our shadow correction method uses gamma correction (0.6) on the LAB color space's lightness channel. This approach effectively brightens shadows while preserving color balance. However, it introduces quantization artifacts due to the non-linear transformation of discrete pixel values. The resulting histograms show noise-like patterns from uneven intensity redistribution. This exemplifies a common image processing trade-off: improved shadow visibility versus introduced artifacts. Potential improvements include higher bit-depth processing or adaptive techniques. Overall, the method achieves its primary goal of shadow enhancement, but with some compromises in image quality.

4 Question 4: Enhancing Vibrance

```
1 image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)
2 image_hsv = cv2.cvtColor(image_rgb, cv2.COLOR_RGB2HSV)
3 # Split into Hue, Saturation, and Value channels
4 hue, saturation, value = cv2.split(image_hsv)
5 def intensity_transformation(saturation, a, sigma=70):
6     # Precompute constant values
7     exponent = np.exp(-((saturation - 128) ** 2) / (2 * (sigma
8         ** 2)))
8     # Apply the intensity transformation
```

```

9     transformed_saturation = np.clip(saturation + a * 128 * 
10    exponent, 0, 255)
11    return transformed_saturation
12 a_value=0.6
13 transformed_saturation = intensity_transformation(saturation.
14    astype(np.float32), a=a_value)
15 # Recombine the transformed saturation with the original hue and
16    value
17 transformed_hsv = cv2.merge([hue, transformed_saturation.astype(
18    np.uint8), value])
19 # Convert the HSV image back to RGB
20 transformed_image_rgb = cv2.cvtColor(transformed_hsv, cv2.
21    COLOR_HSV2RGB)

```



Figure 5: Vibrance Enhanced Image

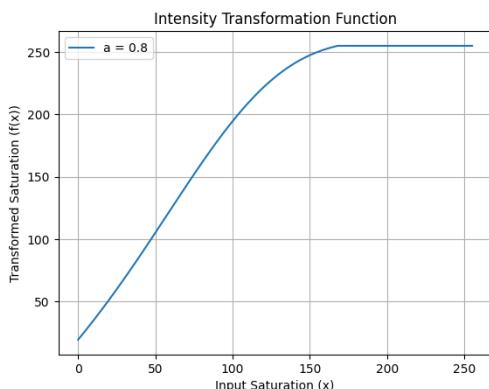


Figure 6: Intensity Transformation Function

Discussion

Our vibrance enhancement method targets the saturation plane in HSV color space. We observed that colorful areas have high saturation values, while colorless areas have low values. We applied an intensity transformation

$$f(x) = \min \left(x + a \cdot 128 \cdot e^{-\frac{(x-128)^2}{2\sigma^2}}, 255 \right)$$

to the saturation channel, with $a = 0.6$ and $\sigma = 70$. This approach amplifies colorful areas while leaving colorless regions unchanged. The exponential term ensures the effect is strongest around mid-range saturation values, resulting in a natural-looking enhancement. We then recombined the transformed saturation with the original hue and value channels to produce the final vibrance-enhanced image.

5 Question 5: Histogram Equalization

```

1 hist, bins = np.histogram(shell_image.flatten(), 256, [0, 256])
2 cdf = hist.cumsum()
3 image_equalized = np.interp(shell_image.flatten(), np.arange(0,
4    256), cdf_normalized)
5 image_equalized = image_equalized.reshape(shell_image.shape)

```

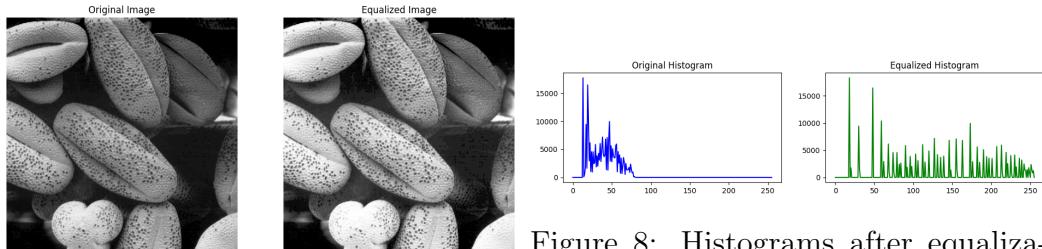


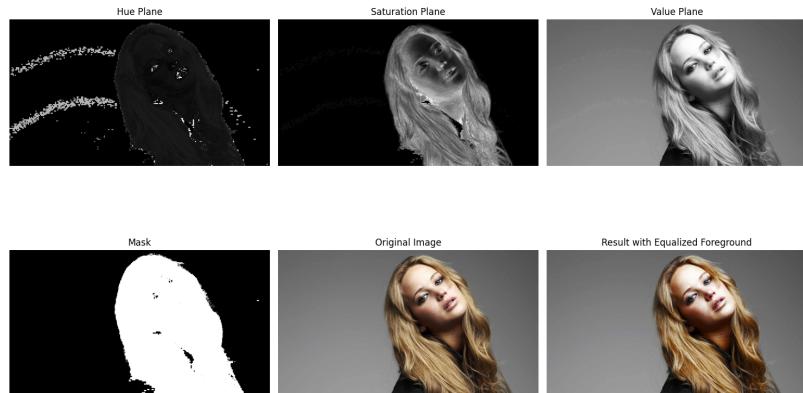
Figure 7: Output image after histogram equalization

Figure 8: Histograms after equalization

Discussion: The histogram equalization applied to the shell image has markedly improved its visual quality. Figure 7 shows enhanced contrast and sharper details in the equalized image, particularly in the shells' textures. Figure 8 demonstrates the histogram transformation, with pixel intensities shifting from a concentrated lower range to a more even distribution. This redistribution unveils previously hidden details in darker areas, though it may slightly increase noise. Overall, the technique effectively enhances the image's dynamic range, resulting in a more informative and visually striking representation of the shell structures.

6 Question 6: Histogram Equalization of Foreground

```
1 jeniffer_image_hsv = cv2.cvtColor(jeniffer_image, cv2.  
    COLOR_BGR2HSV)  
2 hue, saturation, value = cv2.split(jeniffer_image_hsv)  
3 ret, mask = cv2.threshold(saturation, 15, 255, cv2.THRESH_BINARY)  
4 foreground = cv2.bitwise_and(saturation, saturation, mask=mask)  
5 foreground_hist, bins = np.histogram(foreground[foreground > 0],  
    256, [0, 256])  
6 # Compute the CDF (Cumulative Distribution Function)  
7 cdf = foreground_hist.cumsum()  
8 cdf_normalized = (cdf - cdf.min()) * 255 / (cdf.max() - cdf.min())  
9 cdf_normalized = cdf_normalized.astype('uint8')  
10 foreground_equalized = cdf_normalized[foreground]  
11 background = cv2.bitwise_and(saturation, saturation, mask=cv2.  
    bitwise_not(mask))  
12 combined_image = cv2.add(foreground_equalized, background)  
13 hsv_equalized = cv2.merge([hue, combined_image, value])  
14 final_image = cv2.cvtColor(hsv_equalized, cv2.COLOR_HSV2BGR)
```



Approach: This code performs selective histogram equalization on an image's foreground. It converts the image to HSV color space, creates a mask based on saturation to identify the foreground, and applies histogram equalization only to these areas. The process involves calculating the histogram of foreground pixels, deriving a normalized cumulative distribution function (CDF), and using it to remap pixel intensities. The background remains unchanged. Finally, the equalized foreground is combined with the original background and converted back to BGR color space. This method enhances contrast in the foreground while preserving the background, potentially improving the overall image quality and visual appeal.

7 Question 7: Sobel Filtering

```
1 def sobel_filter(image, kernel):
2     [h, w] = np.shape(image) # Get rows and columns of the image
3     output= np.zeros(shape=(h, w))
4     for i in range(1,h-1):
5         for j in range(1, w-1):
6             output[i, j] = np.sum(np.multiply(kernel,image[i-1:i
7 + 2, j-1:j + 2]))
8     return output
9 # Manually apply Sobel filters
10 sobel_manual_x = sobel_filter(image, sobel_x)
11 sobel_manual_y = sobel_filter(image, sobel_y)
12 # Sobel filter property split into two 1D filters
13 kernel_x_row = np.array([[1, 0, -1]]) # 1x3
14 kernel_x_col = np.array([[1], [2], [1]]) # 3x1
15 # Apply separable Sobel filtering
16 sobel_separable_x = cv2.filter2D(image, -1, kernel_x_row)
17 sobel_separable_x = cv2.filter2D(sobel_separable_x, -1,
18     kernel_x_col)
19 # Sobel filter in the Y direction using the same property
20 kernel_y_row = np.array([[1], [0], [-1]]) # 3x1
21 kernel_y_col = np.array([[1, 2, 1]]) # 1x3
22 sobel_separable_y = cv2.filter2D(image, -1, kernel_y_row)
23 sobel_separable_y = cv2.filter2D(sobel_separable_y, -1,
24     kernel_y_col)
```

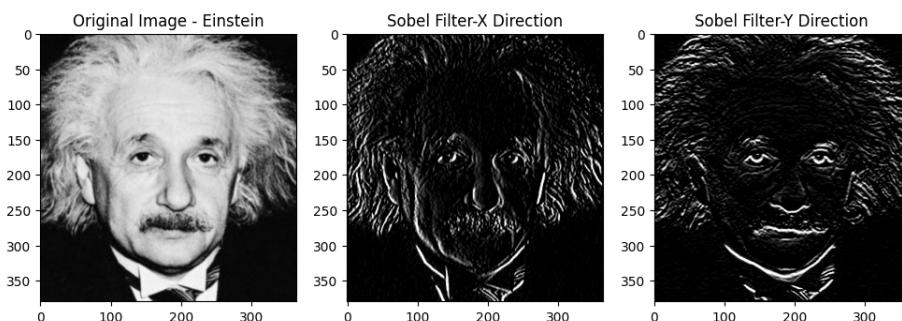


Figure 9: Existing filter2D output

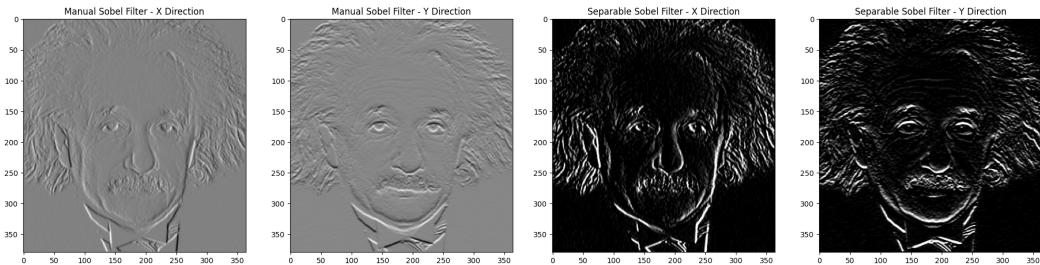


Figure 10: Custom filter Output

Figure 11: Output using Property

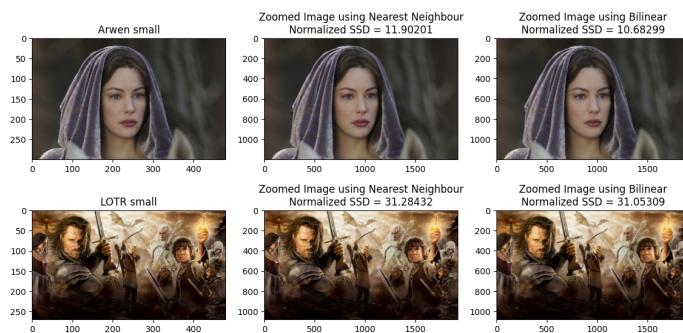
8 Question 8: Image Zooming

Approach: The code implements a `zoom_image` function that resizes images using specified zoom factors and interpolation methods. It handles nearest neighbor and bilinear interpolation, with a fallback for other OpenCV methods. An `ssd` function is also implemented to calculate the normalized Sum of Squared Differences between original and zoomed images, providing a metric to assess zooming quality. This approach allows for flexible image zooming and quantitative quality evaluation.

```

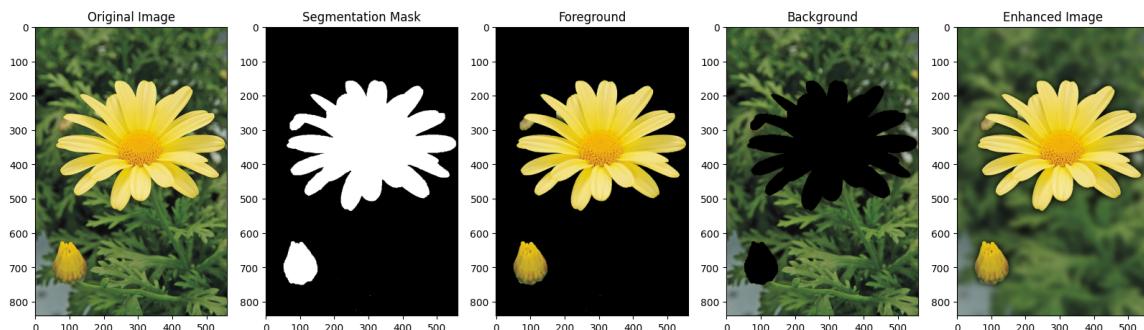
1 def zoom_image(image, zoom_factor, interpolation):
2     M, N, c = image.shape
3     new_M = int(M * zoom_factor)
4     new_N = int(N * zoom_factor)
5     if interpolation == cv2.INTER_NEAREST:
6         zoomed_image = cv2.resize(image, (new_N, new_M),
7             interpolation=cv2.INTER_NEAREST)
7     elif interpolation == cv2.INTER_LINEAR:
8         zoomed_image = cv2.resize(image, (new_N, new_M),
9             interpolation=cv2.INTER_LINEAR)
10    else:
11        try:
12            zoomed_image = cv2.resize(image, (new_N, new_M),
13                interpolation=interpolation)
14        except:
15            raise ValueError('Invalid interpolation method')
16    return zoomed_image
17
# Calculate SSD
def ssd(original_image, zoomed_image):
    return np.sum(np.square(original_image - zoomed_image)) /
    original_image.size

```



9 Question 9: Image Segmentation Using Grab-Cut

```
1 daisy_rgb = cv2.cvtColor(daisy_image, cv2.COLOR_BGR2RGB)
2 M,N,c = daisy_rgb .shape
3 # Define an initial mask and rectangle
4 mask = np.zeros(daisy_rgb .shape[:2], np.uint8)
5 rect = (50, 50, daisy_rgb .shape[1] - 50, daisy_rgb .shape[0] -
50)
6 bkgd_model = np.zeros((1, 65), np.float64)
7 fgd_model = np.zeros((1, 65), np.float64)
8 # Apply grabCut algorithm
9 cv2.grabCut(daisy_rgb , mask, rect, bkgd_model, fgd_model, 5, cv2
    .GC_INIT_WITH_RECT)foreground are marked
10 final_mask = np.where((mask == 2) | (mask == 0), 0, 1).astype('
    uint8')
11 foreground = daisy_rgb * final_mask[:, :, np.newaxis]
12 background = daisy_rgb * (1 - final_mask[:, :, np.newaxis])
13 def apply_background_blur(image, mask, blur_strength=55):
14     blurred = cv2.GaussianBlur(image, (blur_strength,
15         blur_strength), 0)
16     # Use the mask to combine the original foreground with the
17     # blurred background
18     result = np.where(mask[:, :, np.newaxis] == 1, image, blurred)
19     return result
20 # Create enhanced image with blurred background
21 enhanced_image = apply_background_blur(background, final_mask)
22 enhanced_image =cv2.add(foreground, enhanced_image)
```



(c) Reason edge of the flower quite dark in the enhanced image :The dark area around the flower's edge in the enhanced image is caused by the blurring process. When applying the Gaussian blur to the background, pixels near the foreground-background border are averaged with black pixels (zero values) from the foreground mask. This averaging creates a darkening effect just beyond the flower's outline, resulting in the observed darker region in the background close to the flower's edge.

You can find my notebook from [github](#).